

Neural Modeling and Computational Neuroscience

Claudio Gallicchio, Ph.D.



Computational Intelligence & Machine Learning

<http://www.di.unipi.it/groups/ciml>



Dipartimento di Informatica
Università di Pisa - Italy

Additional Info

- 1) Using state-space models for next-step time-series prediction
- 2) Hold out cross validation

LSMs for solving a learning problem

- load the dataset
 - `load laser_dataset.mat`
 - `data = cell2mat(laserTargets)`
- split input & target output
 - `input = data(1:end-1)`
 - `target = data(2:end)`
- divide the data in training set and test set
 - `tr_input = input(1:1500)`
 - `ts_input = input(1500+1:2000)`
 - `tr_target = target(1:1500)`
 - `ts_target = target(1500+1:2000)`

LSMs for solving a learning problem

- Run the liquid using the input of the training set

- you will need to change this line:

```
I=[5*randn(Ne,1);2*randn(Ni,1)]; % thalamic input
```

into something like this:

```
I=[scaling_p1 * tr_input(t) * ones(Ne,1); scaling_p2 * tr_input(t)  
*ones(Ni,1)];
```

- where `scaling_p1` and `scaling_p2` are two hyper-parameters

LSMs for solving a learning problem

- Collect the activations of the liquid for the training samples, i.e.:
 - put the value of the membrane potential of every neuron in the liquid into a column vector, e.g. `state(t)`
 - You can try other choices for the neuronal encoding (e.g. firing rates, or moving averages over the spiking activations)
 - concatenate these columns into a matrix, e.g. `trainStates`
 - analogously, put the target for the training samples into a matrix (a row vector, in this case), e.g. `trainTargets = tr_target`

LSMs for solving a learning problem

- Solve the LMS problem in closed form using, e.g., pseudo-inversion
 - `Wout = trainTargets * pinv(trainStates)`
 - this makes `tr_output = Wout * trainStates` approximate the trainTargets vector in the LMS sense
- You can now compute the MAE on the training set
 - compute the output: `tr_output = Wout * trainStates`
 - compare with the target:
`tr_error = mean(abs(tr_output - trainTargets))`

LSMs for solving a learning problem

- Compute the output for the test set
 - run the liquid on the test set time-steps
now you will have to use `ts_input` instead of `tr_input`
 - as done for the training set, collect the states of the test samples into a matrix, e.g. `testStates`
 - compute the output on the test set:
`ts_output = Wout * testStates`
- Compute the MAE on the test set
 - `ts_error = mean(abs(ts_output - ts_target))`

Tuning the values of the hyper-parameters

- The **performance** of your model **depends on** the values of some crucial **hyper-parameters**, e.g., the number of neurons, the scaling of the input, etc.
- The values of these hyper-parameters should be determined by *model selection*, **without looking at the test set**

Hold out cross validation

- Put the test set aside for now
- Split the training samples into a training set and a validation set (e.g., 70%-30%)
- Train on the training set and evaluate the performance on the validation set for your network, using multiple values of the hyper-parameters
- Choose the set of hyp.-par. values that minimize the MAE on the validation set

Refit

With the chosen values of the hyper-parameters:

- train on the original (complete) training set
- evaluate on the test set